

# Table of contents

Getting started with docker .....	2
Dockerizing MongoDB .....	12
Dockerizing PostgreSQL .....	19
Dockerizing MySQL .....	27
Dockerizing MSSQL .....	35
Dockerizing React, Node.js, Postgres & Nginx <Dev & Prod> .....	44
Restore the old Context Menu in Windows 11 .....	81
Kubernetes Multi Container Deployment   React   Node.js   Postgres   Ingress Nginx   step by step .....	82
What is Kubernetes? .....	89
Kubernetes Cheat Sheet .....	104

# Getting started with docker

## Chapter 1: Introduction to Docker

### 1.1 What is Docker?

Docker is a platform that enables developers to package, deploy, and run applications in containers. Containers include everything needed to run the application, making it portable and consistent across different environments.

### 1.2 Why Use Docker?

- Consistency across development, testing, and production environments
- Isolation and security
- Simplified dependency management
- Efficient resource utilization

## Chapter 2: Installing Docker on Windows

### 2.1 Docker Installation

1. Download Docker Desktop from the Docker website (<https://www.docker.com/products/docker-desktop>).
2. Run the installer and follow the on-screen instructions.
3. After installation, launch Docker Desktop.
4. Ensure Docker is running by opening a terminal and typing:

```
docker --version
```

## Chapter 3: Docker Basics

### 3.1 Docker Architecture

- **Docker Client:** CLI to interact with Docker.
- **Docker Daemon:** Runs on the host machine, manages Docker objects.
- **Docker Images:** Read-only templates to create containers.
- **Docker Containers:** Running instances of Docker images.
- **Docker Registry:** Stores Docker images.

## 3.2 Hello World in Docker

1. Open PowerShell or Command Prompt.
2. Run your first container:

```
docker run hello-world
```

## 3.3 Docker CLI Basics

- List Docker CLI commands:

```
docker
```

- Get help on a command:

```
docker <command> --help
```

# Chapter 4: Working with Docker Images

## 4.1 Pulling Images

- Pull an image from Docker Hub:

```
docker pull node
```

## 4.2 Listing Images

- List all images on your system:

```
docker images
```

## 4.3 Removing Images

- Remove an image:

```
docker rmi <image_id>
```

# Chapter 5: Docker Containers

## 5.1 Running Containers

- Run a Node.js container interactively:

```
docker run -it node /bin/bash
```

- Run a container in the background:

```
docker run -d node
```

## 5.2 Listing Containers

- List all running containers:

```
docker ps
```

- List all containers (including stopped):

```
docker ps -a
```

## 5.3 Stopping Containers

- Stop a running container:

```
docker stop <container_id>
```

## 5.4 Removing Containers

- Remove a container:

```
docker rm <container_id>
```

# Chapter 6: Dockerfile

## 6.1 Introduction to Dockerfile

A Dockerfile is a text document that contains instructions for building a Docker image.

## 6.2 Creating a Dockerfile for Node.js

1. Create a directory for your Node.js application, e.g., `my-node-app`.
2. Inside this directory, create a file named `Dockerfile`.
3. Add the following content:

```
# Use an official Node.js runtime as a parent image
FROM node:14

# Set the working directory in the container
WORKDIR /usr/src/app

# Copy package.json and package-lock.json
COPY package*.json ./

# Install dependencies
RUN npm install

# Copy the rest of the application code
COPY . .

# Expose port 3000
EXPOSE 3000
```

```
# Command to run the application  
CMD ["node", "index.js"]
```

## 6.3 Building an Image

- Build an image from the Dockerfile:

```
docker build -t my-node-app .
```

## 6.4 Running Your Image

- Run the image as a container:

```
docker run -p 3000:3000 my-node-app
```

# Chapter 7: Docker Volumes

## 7.1 Introduction to Volumes

Volumes are used to persist data generated by and used by Docker containers.

## 7.2 Creating Volumes

- Create a volume:

```
docker volume create my-volume
```

## 7.3 Using Volumes

- Use a volume in a container:

```
docker run -d -v my-volume:/usr/src/app my-node-app
```

# Chapter 8: Docker Compose

## 8.1 Introduction to Docker Compose

Docker Compose is a tool for defining and running multi-container Docker applications.

## 8.2 Creating a docker-compose.yml for Node.js and React

1. In your project directory, create a file named `docker-compose.yml`.
2. Add the following content:

```
version: '3'

services:
  web:
    image: my-node-app
    build: .
    ports:
      - "3000:3000"
    volumes:
      - ./usr/src/app
    environment:
      - NODE_ENV=development
  client:
    image: node:14
    working_dir: /usr/src/app
    volumes:
      - ./client:/usr/src/app
    command: npm start
    ports:
      - "3001:3001"
```

## 8.3 Running Docker Compose

- Start your application:

```
docker-compose up
```

- Stop your application:

```
docker-compose down
```

## Chapter 9: Docker Networking

### 9.1 Introduction to Docker Networking

Docker provides a networking model to allow containers to communicate with each other and with non-Docker workloads.

### 9.2 Listing Networks

- List all Docker networks:

```
docker network ls
```

### 9.3 Creating a Network

- Create a custom network:

```
docker network create my-network
```

### 9.4 Connecting Containers to a Network

- Connect a container to a network:

```
docker network connect my-network <container_id>
```

### 9.5 Disconnecting Containers from a Network

- Disconnect a container from a network:

```
docker network disconnect my-network <container_id>
```

## Chapter 10: Docker Swarm



## 10.1 Introduction to Docker Swarm

Docker Swarm is a container orchestration tool that allows you to manage a cluster of Docker nodes.

## 10.2 Initializing a Swarm

- Initialize a swarm:

```
docker swarm init
```

## 10.3 Joining a Swarm

- Get the join command from the manager node and run it on the worker node:

```
docker swarm join --token <token> <manager_ip>:2377
```

## 10.4 Deploying a Service

- Deploy a service in the swarm:

```
docker service create --name my-web-service -p 3000:3000 my-node-app
```

## 10.5 Listing Services

- List all services in the swarm:

```
docker service ls
```

## 10.6 Removing a Service

- Remove a service:

```
docker service rm my-web-service
```

# Chapter 11: Docker Best Practices

## 11.1 Writing Efficient Dockerfiles

- Use official images as a base.
- Minimize the number of layers.
- Use multi-stage builds for optimized images.

## 11.2 Managing Secrets

- Use Docker secrets to manage sensitive data:

```
echo "my_secret_password" | docker secret create my_secret -
```

## 11.3 Security Practices

- Run containers as a non-root user.
- Keep the host and Docker up to date.

# Chapter 12: Advanced Topics

## 12.1 Docker with Kubernetes

- Install and configure Kubernetes.
- Deploy Docker containers using Kubernetes.

## 12.2 CI/CD with Docker

- Use Docker in your CI/CD pipeline.
- Example with Jenkins:

```
pipeline {
  agent any
  stages {
    stage('Build') {
      steps {
        script {
          dockerImage = docker.build("my-node-app")
        }
      }
    }
    stage('Test') {
      steps {
        script {
          dockerImage.inside {
            sh 'npm test'
          }
        }
      }
    }
    stage('Deploy') {
      steps {
        script {
          dockerImage.push('my-repo/my-node-app')
        }
      }
    }
  }
}
```

# Dockerizing MongoDB

## Chapter 1: Introduction

This guide will walk you through the steps of Dockerizing MongoDB, setting up a user with specific credentials, persisting data using Docker volumes, and performing basic CRUD (Create, Read, Update, Delete) operations. We will cover everything from pulling the MongoDB image to running queries.

## Chapter 2: Setting Up Docker

### 2.1 Installing Docker on Windows

1. Download Docker Desktop from the Docker website (<https://www.docker.com/products/docker-desktop>).
2. Run the installer and follow the instructions.
3. After installation, start Docker Desktop.
4. Verify the installation:

```
docker --version
```

## Chapter 3: Dockerizing MongoDB

### 3.1 Pulling the MongoDB Image

- Open a command prompt and run the following command to pull the official MongoDB image:

```
docker pull mongo
```

### 3.2 Creating a Docker Volume

- Create a volume to persist MongoDB data:

```
docker volume create mongodb-data
```

### 3.3 Running MongoDB Container with Authentication and Persistent Storage

- Run the MongoDB container with environment variables to set the username and password, and use the volume for data persistence:

```
docker run --name mongodb_container -d -p 27017:27017 -e  
MONGO_INITDB_ROOT_USERNAME=admin -e MONGO_INITDB_ROOT_PASSWORD=pass -v  
mongodb-data:/data/db mongo
```

- Verify the container is running:

```
docker ps
```

## Chapter 4: Connecting to MongoDB

### 4.1 Using MongoDB Shell

You have two options to connect to MongoDB shell:

#### Option 1: Direct Command

- Start the MongoDB shell directly:

```
docker exec -it mongodb_container bash -c 'mongosh -u admin -p pass --  
authenticationDatabase admin'
```

#### Option 2: Entering the Container First

1. Start a bash shell inside the running MongoDB container:

```
docker exec -it mongodb_container /bin/bash
```

2. Once inside the container, start the MongoDB shell:

```
mongosh -u admin -p pass --authenticationDatabase admin
```

3. You should now be in the MongoDB shell:

```
>
```

4. List all databases:

```
> show databases
```

## Chapter 5: CRUD Operations

### 5.1 Creating a Database and Collection

1. Create a new database called `mydatabase`:

```
use mydatabase
```

2. Create a new collection called `mycollection`:

```
db.createCollection("mycollection")
```

### 5.2 Create (Insert) Documents

1. Insert a single document into `mycollection`:

```
db.mycollection.insertOne({ name: "John Doe", age: 30, occupation: "Engineer" })
```

2. Insert multiple documents:

```
db.mycollection.insertMany([  
  { name: "Jane Doe", age: 25, occupation: "Teacher" },
```

```
{ name: "Steve Smith", age: 40, occupation: "Chef" }  
])
```

## 5.3 Read (Query) Documents

1. Find one document:

```
db.mycollection.findOne({ name: "John Doe" })
```

2. Find all documents:

```
db.mycollection.find()
```

3. Find documents with a condition:

```
db.mycollection.find({ age: { $gt: 30 } })
```

## 5.4 Update Documents

1. Update a single document:

```
db.mycollection.updateOne({ name: "John Doe" }, { $set: { age: 31 } })
```

2. Update multiple documents:

```
db.mycollection.updateMany({ occupation: "Chef" }, { $set: {  
  occupation: "Head Chef" } })
```

## 5.5 Delete Documents

1. Delete a single document:

```
db.mycollection.deleteOne({ name: "John Doe" })
```

2. Delete multiple documents:

```
db.mycollection.deleteMany({ age: { $lt: 30 } })
```

## Chapter 6: Accessing MongoDB from an Application

### 6.1 Using MongoDB with Node.js

1. Install Node.js from the official website (<https://nodejs.org/>).
2. Create a new project directory and navigate into it:

```
mkdir my-mongo-app  
cd my-mongo-app
```

3. Initialize a new Node.js project:

```
npm init -y
```

4. Install the MongoDB driver:

```
npm install mongodb
```

5. Create an `index.js` file and add the following code:

```
const { MongoClient } = require('mongodb');  
  
async function main() {  
  const uri = "mongodb://admin:pass@localhost:27017/?  
authSource=admin";  
  const client = new MongoClient(uri);  
  
  try {  
    await client.connect();  
  
    const database = client.db('mydatabase');  
    const collection = database.collection('mycollection');  
  
    // Insert a document
```



```

    const insertResult = await collection.insertOne({ name: "Alice",
age: 28, occupation: "Designer" });
    console.log('Inserted document:', insertResult.insertedId);

    // Find a document
    const findResult = await collection.findOne({ name: "Alice" });
    console.log('Found document:', findResult);

    // Update a document
    const updateResult = await collection.updateOne({ name: "Alice" },
{ $set: { age: 29 } });
    console.log('Updated document:', updateResult.modifiedCount);

    // Delete a document
    const deleteResult = await collection.deleteOne({ name: "Alice"
});
    console.log('Deleted document:', deleteResult.deletedCount);
  } finally {
    await client.close();
  }
}

main().catch(console.error);

```

6. Run the application:

```
node index.js
```

## Chapter 7: Cleaning Up

### 7.1 Stopping and Removing the MongoDB Container

- Stop the container:

```
docker stop mongodb_container
```

- Remove the container:

```
docker rm mongodb_container
```

## 7.2 Removing the MongoDB Image

- Remove the MongoDB image:

```
docker rmi mongo
```

## 7.3 Removing the Docker Volume

- Remove the Docker volume:

```
docker volume rm mongodb-data
```

# Dockerizing PostgreSQL

## Chapter 1: Introduction

This guide will walk you through the steps of Dockerizing PostgreSQL, setting up a user with specific credentials, persisting data using Docker volumes, and performing basic CRUD (Create, Read, Update, Delete) operations. We will cover everything from pulling the PostgreSQL image to running queries.

## Chapter 2: Setting Up Docker

### 2.1 Installing Docker on Windows

1. Download Docker Desktop from the Docker website (<https://www.docker.com/products/docker-desktop>).
2. Run the installer and follow the instructions.
3. After installation, start Docker Desktop.
4. Verify the installation:

```
docker --version
```

## Chapter 3: Dockerizing PostgreSQL

### 3.1 Pulling the PostgreSQL Image

- Open a command prompt and run the following command to pull the official PostgreSQL image:

```
docker pull postgres
```

### 3.2 Creating a Docker Volume

- Create a volume to persist PostgreSQL data:

```
docker volume create postgres-data
```

### 3.3 Running PostgreSQL Container with Authentication and Persistent Storage

- Run the PostgreSQL container with environment variables to set the username and password, and use the volume for data persistence:

```
docker run --name postgres_container -d -p 5432:5432 -e  
POSTGRES_USER=admin -e POSTGRES_PASSWORD=pass -v postgres-  
data:/var/lib/postgresql/data postgres
```

- Verify the container is running:

```
docker ps
```

## Chapter 4: Connecting to PostgreSQL

### 4.1 Using psql Shell

You have two options to connect to the PostgreSQL shell:

#### Option 1: Direct Command

- Start the PostgreSQL shell directly:

```
docker exec -it postgres_container psql -U admin
```

#### Option 2: Entering the Container First

1. Start a bash shell inside the running PostgreSQL container:

```
docker exec -it postgres_container /bin/bash
```

2. Once inside the container, start the PostgreSQL shell:

```
psql -U admin
```

3. You should now be in the PostgreSQL shell:

```
admin=#
```

4. In the PostgreSQL shell (psql), you can list all databases using the following command:

- List All Databases

```
\l
```

- or

```
\list
```

5. Additional Useful Commands

- List all tables in the current database:

```
\dt
```

- List all schemas in the current database:

```
\dn
```

- List all users:

```
\du
```

- List all indexes:

```
\di
```

## Chapter 5: CRUD Operations

## 5.1 Creating a Database and Table

1. Create a new database called `mydatabase`:

```
CREATE DATABASE mydatabase;
```

2. Connect to the new database:

```
\c mydatabase
```

3. Create a new table called `mytable`:

```
CREATE TABLE mytable (  
    id SERIAL PRIMARY KEY,  
    name VARCHAR(100),  
    age INT,  
    occupation VARCHAR(100)  
);
```

## 5.2 Create (Insert) Records

1. Insert a single record into `mytable`:

```
INSERT INTO mytable (name, age, occupation) VALUES ('John Doe', 30,  
'Engineer');
```

2. Insert multiple records:

```
INSERT INTO mytable (name, age, occupation) VALUES  
( 'Jane Doe', 25, 'Teacher'),  
( 'Steve Smith', 40, 'Chef');
```

## 5.3 Read (Query) Records

1. Select one record:

```
SELECT * FROM mytable WHERE name = 'John Doe';
```

2. Select all records:

```
SELECT * FROM mytable;
```

3. Select records with a condition:

```
SELECT * FROM mytable WHERE age > 30;
```

## 5.4 Update Records

1. Update a single record:

```
UPDATE mytable SET age = 31 WHERE name = 'John Doe';
```

2. Update multiple records:

```
UPDATE mytable SET occupation = 'Head Chef' WHERE occupation = 'Chef';
```

## 5.5 Delete Records

1. Delete a single record:

```
DELETE FROM mytable WHERE name = 'John Doe';
```

2. Delete multiple records:

```
DELETE FROM mytable WHERE age < 30;
```

# Chapter 6: Accessing PostgreSQL from an Application

## 6.1 Using PostgreSQL with Node.js

1. Install Node.js from the official website (<https://nodejs.org/>).

2. Create a new project directory and navigate into it:

```
mkdir my-postgres-app  
cd my-postgres-app
```

3. Initialize a new Node.js project:

```
npm init -y
```

4. Install the `pg` package:

```
npm install pg
```

5. Create an `index.js` file and add the following code:

```
const { Client } = require('pg');  
  
async function main() {  
  const client = new Client({  
    user: 'admin',  
    host: 'localhost',  
    database: 'mydatabase',  
    password: 'pass',  
    port: 5432,  
  });  
  
  await client.connect();  
  
  try {  
    // Insert a record  
    const insertResult = await client.query("INSERT INTO mytable  
(name, age, occupation) VALUES ('Alice', 28, 'Designer') RETURNING  
id");  
    console.log('Inserted record ID:', insertResult.rows[0].id);  
  
    // Select a record  
    const selectResult = await client.query("SELECT * FROM mytable
```



```

WHERE name = 'Alice'");
    console.log('Selected record:', selectResult.rows[0]);

    // Update a record
    const updateResult = await client.query("UPDATE mytable SET age =
29 WHERE name = 'Alice'");
    console.log('Updated record count:', updateResult.rowCount);

    // Delete a record
    const deleteResult = await client.query("DELETE FROM mytable WHERE
name = 'Alice'");
    console.log('Deleted record count:', deleteResult.rowCount);
  } finally {
    await client.end();
  }
}

main().catch(console.error);

```

6. Run the application:

```
node index.js
```

## Chapter 7: Cleaning Up

### 7.1 Stopping and Removing the PostgreSQL Container

- Stop the container:

```
docker stop postgres_container
```

- Remove the container:

```
docker rm postgres_container
```

### 7.2 Removing the PostgreSQL Image

- Remove the PostgreSQL image:

```
docker rmi postgres
```

## 7.3 Removing the Docker Volume

- Remove the Docker volume:

```
docker volume rm postgres-data
```

# Dockerizing MySQL

## Chapter 1: Introduction

This guide will walk you through the steps of Dockerizing MySQL, setting up a user with specific credentials, persisting data using Docker volumes, and performing basic CRUD (Create, Read, Update, Delete) operations. We will cover everything from pulling the MySQL image to running queries.

## Chapter 2: Setting Up Docker

### 2.1 Installing Docker on Windows

1. Download Docker Desktop from the Docker website (<https://www.docker.com/products/docker-desktop>).
2. Run the installer and follow the instructions.
3. After installation, start Docker Desktop.
4. Verify the installation:

```
docker --version
```

## Chapter 3: Dockerizing MySQL

### 3.1 Pulling the MySQL Image

- Open a command prompt and run the following command to pull the official MySQL image:

```
docker pull mysql
```

### 3.2 Creating a Docker Volume

- Create a volume to persist MySQL data:

```
docker volume create mysql-data
```

### 3.3 Running MySQL Container with Authentication and Persistent Storage

- Run the MySQL container with environment variables to set the username and password, and use the volume for data persistence:

```
docker run --name mysql_container -d -p 3306:3306 -e  
MYSQL_ROOT_PASSWORD=pass -e MYSQL_USER=admin -e MYSQL_PASSWORD=pass -e  
MYSQL_DATABASE=mydatabase -v mysql-data:/var/lib/mysql mysql
```

- Verify the container is running:

```
docker ps
```

## Chapter 4: Connecting to MySQL

### 4.1 Using MySQL Shell

You have two options to connect to the MySQL shell:

#### Option 1: Direct Command

- Start the MySQL shell directly:

```
docker exec -it mysql_container mysql -u admin -p
```

- you will be prompted to Enter password:

```
Enter password: pass
```

#### Option 2: Entering the Container First

1. Start a bash shell inside the running MySQL container:

```
docker exec -it mysql_container /bin/bash
```

2. Once inside the container, start the MySQL shell:

```
mysql -u admin -p
```

3. Enter the password when prompted (pass in this example).

4. You should now be in the MySQL shell:

```
mysql>
```

5. In the MySQL shell, you can list all databases using the following command:

- List All Databases:

```
SHOW DATABASES;
```

## Chapter 5: CRUD Operations

### 5.1 Creating a Database and Table

1. Create a new database called mydatabase:

```
CREATE DATABASE mydatabase;
```

2. Select the new database:

```
USE mydatabase;
```

3. Create a new table called mytable:

```
CREATE TABLE mytable (  
    id INT AUTO_INCREMENT PRIMARY KEY,  
    name VARCHAR(100),  
    age INT,  
    occupation VARCHAR(100)  
);
```

## 5.2 Create (Insert) Records

1. Insert a single record into mytable:

```
INSERT INTO mytable (name, age, occupation) VALUES ('John Doe', 30, 'Engineer');
```

2. Insert multiple records:

```
INSERT INTO mytable (name, age, occupation) VALUES ('Jane Doe', 25, 'Teacher'), ('Steve Smith', 40, 'Chef');
```

## 5.3 Read (Query) Records

1. Select one record:

```
SELECT * FROM mytable WHERE name = 'John Doe';
```

2. Select all records:

```
SELECT * FROM mytable;
```

3. Select records with a condition:

```
SELECT * FROM mytable WHERE age > 30;
```

## 5.4 Update Records

1. Update a single record:

```
UPDATE mytable SET age = 31 WHERE name = 'John Doe';
```

2. Update multiple records:

```
UPDATE mytable SET occupation = 'Head Chef' WHERE occupation = 'Chef';
```

## 5.5 Delete Records

1. Delete a single record:

```
DELETE FROM mytable WHERE name = 'John Doe';
```

2. Delete multiple records:

```
DELETE FROM mytable WHERE age < 30;
```

## 5.6 Additional Useful Commands

- List all tables in the current database:

```
SHOW TABLES;
```

- Describe the structure of a table:

```
DESCRIBE mytable;
```

- List all users:

```
SELECT User, Host FROM mysql.user;
```

# Chapter 6: Accessing MySQL from an Application

## 6.1 Using MySQL with Node.js

1. Install Node.js from the official website (<https://nodejs.org/>).
2. Create a new project directory and navigate into it:

```
mkdir my-mysql-app
```

```
cd my-mysql-app
```

3. Initialize a new Node.js project:

```
npm init -y
```

4. Install the `mysql` package:

```
npm install mysql2
```

5. Create an `index.js` file and add the following code:

```
const mysql = require('mysql2/promise');

async function main() {
  const connection = await mysql.createConnection({
    host: 'localhost',
    user: 'admin',
    password: 'pass',
    database: 'mydatabase'
  });

  try {
    console.log('connected as id ' + connection.threadId);

    // Insert a record
    const [insertResults] = await connection.execute(
      "INSERT INTO mytable (name, age, occupation) VALUES (?, ?, ?)",
      ['Alice', 28, 'Designer']
    );
    console.log('Inserted record ID:', insertResults.insertId);

    // Select a record
    const [selectResults] = await connection.execute(
      "SELECT * FROM mytable WHERE name = ?",
      ['Alice']
    );
  } catch (error) {
    console.error('Error:', error);
  }
}
```



```

    );
    console.log('Selected record:', selectResults[0]);

    // Update a record
    const [updateResults] = await connection.execute(
        "UPDATE mytable SET age = ? WHERE name = ?",
        [29, 'Alice']
    );
    console.log('Updated record count:', updateResults.affectedRows);

    // Delete a record
    const [deleteResults] = await connection.execute(
        "DELETE FROM mytable WHERE name = ?",
        ['Alice']
    );
    console.log('Deleted record count:', deleteResults.affectedRows);
} catch (err) {
    console.error('error:', err.stack);
} finally {
    await connection.end();
}
}

main();

```

6. Run the application:

```
node index.js
```

## Chapter 7: Cleaning Up

### 7.1 Stopping and Removing the MySQL Container

- Stop the container:

```
docker stop mysql_container
```

- Remove the container:

```
docker rm mysql_container
```

## 7.2 Removing the MySQL Image

- Remove the MySQL image:

```
docker rmi mysql
```

## 7.3 Removing the Docker Volume

- Remove the Docker volume:

```
docker volume rm mysql-data
```

# Dockerizing MSSQL

## Chapter 1: Introduction

This guide will walk you through the steps of Dockerizing Microsoft SQL Server (MSSQL), setting up a user with specific credentials, persisting data using Docker volumes, and performing basic CRUD (Create, Read, Update, Delete) operations on a Windows system. We will cover everything from pulling the MSSQL image to running queries.

## Chapter 2: Setting Up Docker

### 2.1 Installing Docker on Windows

1. Download Docker Desktop from the Docker website (<https://www.docker.com/products/docker-desktop>).
2. Run the installer and follow the instructions.
3. After installation, start Docker Desktop.
4. Verify the installation:

```
docker --version
```

## Chapter 3: Dockerizing MSSQL Server

### 3.1 Pulling the MSSQL Server Image

- Open a command prompt or PowerShell and run the following command to pull the official MSSQL Server image:

```
docker pull mcr.microsoft.com/mssql/server
```

### 3.2 Creating a Docker Volume

- Create a volume to persist MSSQL Server data:

```
docker volume create mssql-data
```

### 3.3 Running MSSQL Server Container with Authentication and Persistent Storage

- Run the MSSQL Server container with environment variables to set the SA password and use the volume for data persistence:

```
docker run -e "ACCEPT_EULA=Y" -e "SA_PASSWORD=yourStrong(!)Password" -p 1433:1433 --name mssql_container -v mssql-data:/var/opt/mssql -d mcr.microsoft.com/mssql/server
```

- Verify the container is running:

```
docker ps
```

## Chapter 4: Installing MSSQL Command Line Tools

### 4.1 Download and Install MSSQL Tools

1. Download the Microsoft ODBC Driver 17 for SQL Server from the Microsoft website (<https://docs.microsoft.com/en-us/sql/connect/odbc/download-odbc-driver-for-sql-server>).
2. Install the ODBC driver by running the downloaded installer.
3. Download the SQL Server Command Line Tools (sqlcmd and bcp) from the Microsoft website (<https://docs.microsoft.com/en-us/sql/tools/sqlcmd-utility>).
4. Install the SQL Server Command Line Tools by running the downloaded installer.

## Chapter 5: Connecting to MSSQL Server

### 5.1 Using MSSQL Server Command Line Tools

1. Open Command Prompt or PowerShell.
2. Connect to the MSSQL Server using sqlcmd:

```
sqlcmd -S localhost -U SA -P "yourStrong(!)Password"
```

3. You should now be in the MSSQL command line:

```
1>
```

4. In the MSSQL command line, you can list all databases using the following command:

```
SELECT name FROM sys.databases;  
GO
```

## Example Session

- Here's how an example session might look:

```
C:\> sqlcmd -S localhost -U SA -P "yourStrong(!)Password"  
1> SELECT name FROM sys.databases;  
2> GO  
name  
-----  
-----  
master  
tempdb  
model  
msdb  
mydatabase  
  
(5 rows affected)  
1>
```

## Chapter 6: CRUD Operations

### 6.1 Creating a Database and Table

1. Create a new database called `mydatabase`:

```
CREATE DATABASE mydatabase;  
GO
```

2. Use the new database:

```
USE mydatabase;  
GO
```

3. Create a new table called mytable:

```
CREATE TABLE mytable (  
    id INT PRIMARY KEY IDENTITY(1,1),  
    name NVARCHAR(100),  
    age INT,  
    occupation NVARCHAR(100)  
);  
GO
```

## 6.2 Create (Insert) Records

1. Insert a single record into mytable:

```
INSERT INTO mytable (name, age, occupation) VALUES ('John Doe', 30,  
'Engineer');  
GO
```

2. Insert multiple records:

```
INSERT INTO mytable (name, age, occupation) VALUES  
('Jane Doe', 25, 'Teacher'),  
('Steve Smith', 40, 'Chef');  
GO
```

## 6.3 Read (Query) Records

1. Select one record:

```
SELECT * FROM mytable WHERE name = 'John Doe';  
GO
```

2. Select all records:

```
SELECT * FROM mytable;  
GO
```

3. Select records with a condition:

```
SELECT * FROM mytable WHERE age > 30;  
GO
```

## 6.4 Update Records

1. Update a single record:

```
UPDATE mytable SET age = 31 WHERE name = 'John Doe';  
GO
```

2. Update multiple records:

```
UPDATE mytable SET occupation = 'Head Chef' WHERE occupation = 'Chef';  
GO
```

## 6.5 Delete Records

1. Delete a single record:

```
DELETE FROM mytable WHERE name = 'John Doe';  
GO
```

2. Delete multiple records:

```
DELETE FROM mytable WHERE age < 30;
```

```
GO
```

## 6.6 Additional Useful Commands

- List all tables in the current database:

```
SELECT * FROM sys.Tables;  
GO
```

- Describe the structure of a table:

```
sp_help mytable;  
GO
```

- List all users:

```
SELECT name FROM sys.sql_logins;  
GO
```

# Chapter 7: Accessing MSSQL Server from an Application

## 7.1 Using MSSQL Server with Node.js

1. Install Node.js from the official website (<https://nodejs.org/>).
2. Create a new project directory and navigate into it:

```
mkdir my-mssql-app  
cd my-mssql-app
```

3. Initialize a new Node.js project:

```
npm init -y
```

4. Install the `mssql` package:



```
npm install mssql
```

5. Create an `index.js` file and add the following code:

```
const sql = require('mssql');

const config = {
  user: 'sa',
  password: 'yourStrong(!)Password',
  server: 'localhost',
  database: 'mydatabase',
  options: {
    encrypt: true, // Use encryption
    trustServerCertificate: true // For self-signed certificate
  }
};

async function main() {
  try {
    let pool = await sql.connect(config);

    // Insert a record
    let insertResult = await pool.request()
      .query("INSERT INTO mytable (name, age, occupation) VALUES ('Alice', 28, 'Designer')");
    console.log('Inserted record:', insertResult);

    // Insert a many
    let insertResult = await pool.request()
      .query(`INSERT INTO mytable (name, age, occupation) VALUES ('jane doe', 30, 'Designer'), ('kyle Smith', 40, 'Chef')`);
    console.log('Inserted record:', insertResult);

    // Select a record
    let selectResult = await pool.request()
      .query("SELECT * FROM mytable WHERE name = 'Alice'");
```

```

    console.log('Selected record:', selectResult.recordset);

    // Update a record
    let updateResult = await pool.request()
        .query("UPDATE mytable SET age = 29 WHERE name = 'Alice'");
    console.log('Updated record:', updateResult);

    // Delete a record
    let deleteResult = await pool.request()
        .query("DELETE FROM mytable WHERE name = 'Alice'");
    console.log('Deleted record:', deleteResult);

    } catch (err) {
        console.error('SQL error', err);
    }
}

main();

```

6. Run the application:

```
node index.js
```

## Chapter 8: Cleaning Up

### 8.1 Stopping and Removing the MSSQL Server Container

- Stop the container:

```
docker stop mssql_container
``
```

- Remove the container:

```
docker rm mssql_container
```

## 8.2 Removing the MSSQL Server Image

- Remove the MSSQL Server image:

```
docker rmi mcr.microsoft.com/mssql/server
```

## 8.3 Removing the Docker Volume

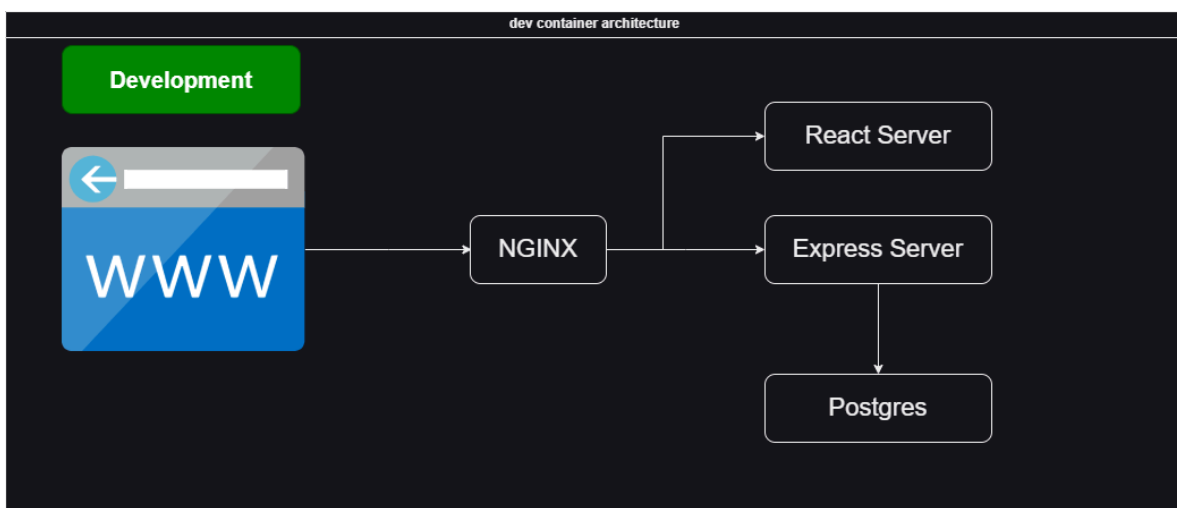
- Remove the Docker volume:

```
docker volume rm mssql-data
```

# Dockerizing React, Node.js, Postgres & Nginx <Dev & Prod>

This guide provides a step-by-step approach to Dockerizing a React application using Vite, with Node.js, Postgres, and Nginx, for both development and production environments.

## Development Mode



dev.png

In your development environment, the architecture works as follows:

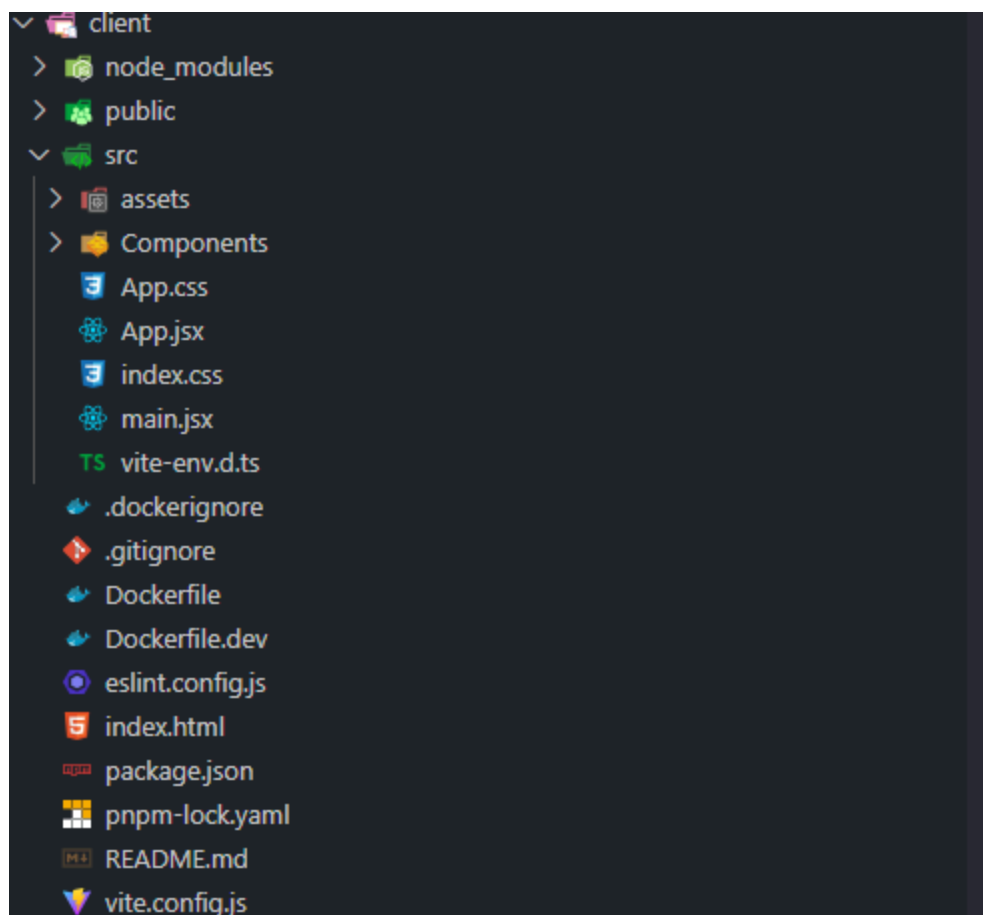
1. **Nginx** acts as a reverse proxy, routing incoming HTTP requests to either the **React server** or the **Express server** based on the request path.
2. **React Server** handles all requests for the frontend (i.e., your React application). When you develop, changes here are served by the React development server, often with hot-reloading.
3. **Express Server** processes backend API requests. These requests might involve querying or updating data in the **Postgres** database.
4. **Nginx** forwards requests:
  - Requests to the base URL (e.g., `http://localhost/`) are sent to the React server.

- Requests to paths starting with `/api` are forwarded to the Express server, which may interact with the Postgres database.

This setup allows you to work on both the frontend and backend in a unified environment, ensuring that you can test how they interact together while still enjoying the benefits of a development-focused workflow like hot-reloading and quick feedback.

## React Folder & File structure

### 1. Client: React + Vite



react-structure.png

- Create a Vite React project:

```
npm create vite@latest my-react-app -- --template react
cd my-react-app
```

```
npm install
```

- **Install necessary dependencies:** You'll need to install `react-router-dom` and `axios` for routing and HTTP requests.

```
npm install react-router-dom@4.3.1 axios
```

## 2. Create the Project Structure

Your project should be structured as follows:

```
my-react-app/  
├─ node_modules/  
├─ public/  
├─ src/  
│   ├─ assets/  
│   ├─ Components/  
│   │   ├─ MainComponent.jsx  
│   │   ├─ MainComponent.css  
│   │   └─ OtherPage.jsx  
│   ├─ App.css  
│   ├─ App.jsx  
│   ├─ index.css  
│   ├─ main.jsx  
│   └─ vite-env.d.ts  
├─ .dockerignore  
├─ .gitignore  
├─ Dockerfile.dev  
├─ Dockerfile  
├─ eslint.config.js  
├─ index.html  
├─ package.json  
├─ pnpm-lock.yaml  
├─ README.md  
└─ vite.config.js
```

## 3. Add the Required Code

- src/App.jsx

```
import './App.css';
import { BrowserRouter as Router, Route, Link } from "react-router-dom";
import MainComponent from './Components/MainComponent';
import OtherPage from './Components/OtherPage';

function App() {
  return (
    <Router>
      <>
        <header className="header">
          <div>This is a multicontainer application</div>
          <Link to="/">Home</Link>
          <Link to="/otherpage">Other page</Link>
        </header>
        <div className="main">
          <Route exact path="/" component={MainComponent} />
          <Route path="/otherpage" component={OtherPage} />
        </div>
      </>
    </Router>
  );
}

export default App;
```

- src/App.css

```
.header {
  background: #eee;
}

.header a {
  margin-left: 20px;
}
```

```
.main {  
  padding: 10px;  
  background: #ccc;  
}
```

- **src/index.css**

```
body {  
  margin: 0;  
  font-family: -apple-system, BlinkMacSystemFont, 'Segoe UI', 'Roboto',  
  'Oxygen',  
    'Ubuntu', 'Cantarell', 'Fira Sans', 'Droid Sans', 'Helvetica Neue',  
    sans-serif;  
  -webkit-font-smoothing: antialiased;  
  -moz-osx-font-smoothing: grayscale;  
}  
  
code {  
  font-family: source-code-pro, Menlo, Monaco, Consolas, 'Courier New',  
    monospace;  
}
```

- **src/main.jsx**

```
import { StrictMode } from 'react';  
import { createRoot } from 'react-dom/client';  
import App from './App.jsx';  
import './index.css';  
  
createRoot(document.getElementById('root')).render(  
  <StrictMode>  
    <App />  
  </StrictMode>  
)
```



```
    </StrictMode>,
  );

```

- **src/Components/MainComponent.jsx**

```
import { useCallback, useState, useEffect } from "react";
import axios from "axios";
import "./MainComponent.css";

function MainComponent() {
  const [values, setValues] = useState([]);
  const [value, setValue] = useState("");

  const getAllNumbers = useCallback(async () => {
    const data = await axios.get("/api/values/all");
    setValues(data.data.rows.map(row => row.number));
  }, []);

  const saveNumber = useCallback(
    async (event) => {
      event.preventDefault();
      await axios.post("/api/values", { value });

      setValue("");
      getAllNumbers();
    },
    [value, getAllNumbers]
  );

  useEffect(() => {
    getAllNumbers();
  }, []);

  return (
    <div>
      <button onClick={getAllNumbers}>Get all numbers</button>

```

```

    <br />
    <span className="title">Values</span>
    <div className="values">
      {values.map((value, index) => (
        <div className="value" key={index}>{value}</div>
      ))}
    </div>
    <form className="form" onSubmit={saveNumber}>
      <label>Enter your value: </label>
      <input
        value={value}
        onChange={event => {
          setValue(event.target.value);
        }}
      />
      <button>Submit</button>
    </form>
  </div>
);
}

export default MainComponent;

```

- src/Components/MainComponent.css

```

.title {
  font-weight: bold;
}

.values {
  margin-top: 20px;
  background: yellow;
}

.value {
  margin-top: 10px;
}

```

```
border-top: 1px dashed black;
}

.form {
  margin-top: 20px;
}
```

- src/Components/OtherPage.jsx

```
import { Link } from "react-router-dom";

function OtherPage() {
  return (
    <div>
      I'm another page!
      <br />
      <br />
      <Link to="/">Go back to home screen</Link>
    </div>
  );
}

export default OtherPage;
```

## 4. Configuration Files

- .dockerignore

```
node_modules
build
.git
.gitignore
Dockerfile
Dockerfile.dev
.dockerignore
```

```
*.log  
.DS_Store  
.git
```

- **.gitignore**

```
# Logs  
logs  
*.log  
npm-debug.log*  
yarn-debug.log*  
yarn-error.log*  
pnpm-debug.log*  
lerna-debug.log*  
  
node_modules  
dist  
dist-ssr  
*.local  
  
# Editor directories and files  
.vscode/*  
!.vscode/extensions.json  
.idea  
.DS_Store  
*.suo  
*.ntvs*  
*.njsproj  
*.sln  
*.sw?
```

#### 4. Adjust the vite.config.js config

- **vite.config.js**

```
import { defineConfig } from 'vite'
import react from '@vitejs/plugin-react'
export default defineConfig({
  plugins: [react()],
  server: {
    host: true,
    port: 5173,
    watch: {
      usePolling: true
    }
  }
})
```

## 5. Add the docker files

- Docker.dev

```
# Use an official Node.js runtime as a parent image
FROM node:20.16.0-alpine
# Set the working directory in the container
WORKDIR /app
# Install dependencies
COPY package.json pnpm-lock.yaml ./
# Copy the rest of the application code
COPY . .
# Expose port 5173 (default Vite port)
EXPOSE 5173
# Command to run the development server
CMD ["npm", "run", "dev"]
```

## Build the Docker Image

Navigate to the root directory of your project where the Dockerfile.dev is located. Run the following command to build the Docker image for development:

```
docker build -t my-react-app-dev -f Dockerfile.dev .
```

This command tells Docker to:

-t client: Tag the image as my-react-app-dev. -f Dockerfile.dev: Use the Dockerfile.dev file for building the image.

## Run the Docker Container

After building the image, you can run the container with the following command:


```
docker run -it --rm -p 5173:5173 client
```

Explanation:

- **-it**: Runs the container interactively.
- **--rm**: Automatically removes the container when it exits. **-p 5173:5173**: Maps port 5173 on your local machine to port 5173 in the container, which is where the Vite development server will be running.

## Node Express Folder & File Structure

Here's a guide to create an Express.js application, dockerizing it for development, and managing environment variables securely using a `keys.js` file.



serverfolderstructure.png

### Step 1: Set Up the Project Directory

```
server/  
├─ node_modules/  
├─ .dockerignore  
├─ Dockerfile.dev  
├─ index.js  
├─ keys.js  
├─ package.json  
├─ pnpm-lock.yaml  
└─ .gitignore
```

Create a new directory for your project:

```
mkdir express-postgres-app  
cd express-postgres-app
```

## Step 2: Initialize the Node.js Project

Run the following command to initialize a new Node.js project:

```
npm init -y
```

This will create a `package.json` file with default settings.

## Step 3: Install Dependencies

Install the required dependencies: Express, PostgreSQL client (pg), CORS, and body-parser:

```
npm install express pg cors body-parser
```

## Step 4: Create the Application Files

### 1. Create index.js

In the root of your project directory, create a file named `index.js` and add the following code:

```
const keys = require("./keys");  
const express = require("express");  
const bodyParser = require("body-parser");
```

```

const cors = require("cors");
const { Pool } = require("pg");

const app = express();

// Middlewares
app.use(cors());
app.use(bodyParser.json());

// Postgres client setup
const pgClient = new Pool({
  user: keys.pgUser,
  host: keys.pgHost,
  database: keys.pgDatabase,
  password: keys.pgPassword,
  port: keys.pgPort
});

pgClient.on("connect", client => {
  client.query("CREATE TABLE IF NOT EXISTS values (number INT)")
    .catch(err => console.log("PG ERROR", err));
});

// Express route definitions
app.get("/", (req, res) => {
  res.send("Hi");
});

// Get all values
app.get("/values/all", async (req, res) => {
  const values = await pgClient.query("SELECT * FROM values");
  res.send(values.rows);
});

// Insert a new value
app.post("/values", async (req, res) => {
  if (!req.body.value) return res.send({ working: false });

```



```

    await pgClient.query("INSERT INTO values(number) VALUES($1)",
[req.body.value]);

    res.send({ working: true });
  });

app.listen(8000, () => {
  console.log("Listening on port 8000");
});

```

## 2. Create keys.js

Create a file named `keys.js` in the root directory with the following content:

```

module.exports = {
  pgUser: process.env.PGUSER,
  pgHost: process.env.PGHOST,
  pgDatabase: process.env.PGDATABASE,
  pgPassword: process.env.PGPASSWORD,
  pgPort: process.env.PGPORT
};

```

This file securely references environment variables for your PostgreSQL connection.

## Step 5: Create Docker Configuration

### 1. Create Dockerfile.dev

Create a file named `Dockerfile.dev` in the root directory with the following content:

```

# Use Alpine Node.js runtime as a parent image
FROM node:20.16.0-alpine

# Set the working directory in the container
WORKDIR /app

# Copy package.json and pnpm-lock.yaml (or package-lock.json)
COPY package*.json pnpm-lock.yaml ./

```

```
# Install dependencies
RUN npm install

# Copy the rest of the application code
COPY . .

# Expose port 8000
EXPOSE 8000

# Command to run the application
CMD ["npm", "run", "dev"]
```

## 2. Create .dockerignore

Create a `.dockerignore` file to ensure unnecessary files aren't copied into the Docker image:

```
node_modules
build
.git
.gitignore
Dockerfile
Dockerfile.dev
.dockerignore
*.log
.DS_Store
.git
```

## Step 6: Define the NPM Scripts

Open your `package.json` and add a script for running the development server:

```
{
  "scripts": {
    "start": "node index.js",
    "dev": "nodemon index.js"
```

```
}  
}
```

Ensure you have nodemon installed globally or as a dev dependency:

```
npm install -g nodemon
```

Or:

```
npm install --save-dev nodemon
```

## Step 7: Running the Application with Docker

### 1. Build the Docker Image

Run the following command to build the Docker image using the development Dockerfile:

```
docker build -t server -f Dockerfile.dev .
```

### 2. Run the Docker Container

Run the Docker container:

```
docker run -it --rm -p 8000:8000 server
```

Explanation:

- `-p 8000:8000`: Maps port 8000 on your machine to port 8000 in the container.

## Step 8: Test the Application

Open your browser or use a tool like Postman to test the endpoints:

- GET request to `http://localhost:8000/` should return "Hi".

## NGINX & File structure

Nginx as a reverse proxy to route traffic between a React frontend (client) and an Express backend (API), using Docker. The configuration involves two files: default.conf (Nginx configuration) and Dockerfile.dev (Docker configuration for Nginx).

- Step 1: The structure looks like this:

```
nginx/  
├─ default.conf      # Nginx configuration file  
└─ Dockerfile.dev    # Dockerfile for building the Nginx image
```

- Step 2: Configure default.conf for Nginx The default.conf file is an Nginx configuration file that routes traffic based on the request paths. Here's what the configuration does:

```
upstream client {  
    server client:5173;  
}  
  
upstream api {  
    server api:8000;  
}  
  
server {  
    listen 80;  
  
    location / {  
        proxy_pass http://client;  
    }  
  
    location /sockjs-node {  
        proxy_pass http://client;  
        proxy_http_version 1.1;  
        proxy_set_header Upgrade $http_upgrade;  
        proxy_set_header Connection "Upgrade";  
    }  
  
    location /api {
```

```
    rewrite /api/(.*) /$1 break;
    proxy_pass http://api;
}
}
```

Explanation:

- Upstream Sections:
  - upstream client: Defines a group of servers for the React frontend running on port 5173.
  - upstream api: Defines a group of servers for the Express backend running on port 8000.
- Server Block:
  - listen 80: Configures Nginx to listen on port 80 (standard HTTP port).
  - location /: Routes root URL requests to the React client.
  - location /sockjs-node: Special handling for WebSocket connections, typically used for hot-reloading in development.
  - location /api: Routes /api requests to the Express backend, rewriting the URL to strip the /api prefix.
- Step 3: Create the Dockerfile.dev for Nginx The Dockerfile.dev will build a Docker image for Nginx that uses your custom configuration. The Dockerfile content is:

```
FROM nginx
COPY ./default.conf /etc/nginx/conf.d/default.conf
```

## Step-by-Step Guide for Configuring Nginx with Docker

This guide walks you through setting up Nginx as a reverse proxy to route traffic between a React frontend (client) and an Express backend (API), using Docker. The configuration

involves two files: `default.conf` (Nginx configuration) and `Dockerfile.dev` (Docker configuration for Nginx).

### Step 1: Understand the Project Structure

From the provided image, the structure looks like this:

```
nginx/
├── default.conf      # Nginx configuration file
└── Dockerfile.dev    # Dockerfile for building the Nginx image
```

### Step 2: Configure `default.conf` for Nginx

The `default.conf` file is an Nginx configuration file that routes traffic based on the request paths. Here's what the configuration does:

```
upstream client {
    server client:5173;
}

upstream api {
    server api:8000;
}

server {
    listen 80;

    location / {
        proxy_pass http://client;
    }

    location /sockjs-node {
        proxy_pass http://client;
        proxy_http_version 1.1;
        proxy_set_header Upgrade $http_upgrade;
        proxy_set_header Connection "Upgrade";
    }

    location /api {
```

```
        rewrite /api/(.*) /$1 break;
        proxy_pass http://api;
    }
}
```

#### Explanation:

- **Upstream Sections:**

- **upstream client:** Defines a group of servers for the React frontend running on port 5173.
- **upstream api:** Defines a group of servers for the Express backend running on port 8000.

- **Server Block:**

- **listen 80:** Configures Nginx to listen on port 80 (standard HTTP port).
- **location /:** Routes root URL requests to the React client.
- **location /sockjs-node:** Special handling for WebSocket connections, typically used for hot-reloading in development.
- **location /api:** Routes `/api` requests to the Express backend, rewriting the URL to strip the `/api` prefix.

#### Step 3: Create the Dockerfile.dev for Nginx

The `Dockerfile.dev` will build a Docker image for Nginx that uses your custom configuration. The Dockerfile content is:

```
FROM nginx
COPY ./default.conf /etc/nginx/conf.d/default.conf
```

#### Explanation:

- **FROM nginx:** This uses the official Nginx image as the base image.

- **COPY ./default.conf /etc/nginx/conf.d/default.conf:** Copies your custom Nginx configuration file (default.conf) into the appropriate directory inside the Nginx container (/etc/nginx/conf.d/default.conf), replacing the default configuration.

## Combined Project Structure with Docker Compose

Here's how to combine your client, server, and nginx components into a single project structure with a docker-compose.yml file that orchestrates the services. This structure is typical for a full-stack application setup using Docker Compose.

### Project Structure

```

my-app/
├── client/                                # React frontend
│   ├── Dockerfile.dev                    # Dockerfile for the React app
│   (development)
│   ├── src/                             # Source code for React
│   ├── public/                          # Public assets for React
│   ├── package.json                     # Package configuration for React
│   └── ...                              # Other React files and folders
├── server/                              # Express backend
│   ├── Dockerfile.dev                    # Dockerfile for the Express app
│   (development)
│   ├── index.js                         # Main entry point for the Express server
│   ├── keys.js                          # Environment configuration file
│   ├── package.json                     # Package configuration for Express
│   └── ...                              # Other server files and folders
├── nginx/                               # Nginx reverse proxy
│   ├── default.conf                     # Nginx configuration file
│   └── Dockerfile.dev                    # Dockerfile for Nginx (development)
└── docker-compose.yml                    # Docker Compose configuration file

```

### docker-compose.yml

Below is the docker-compose.yml file that defines all the services:

```

version: "3.8"
services:
  postgres:

```



```

    image: "postgres:latest"
    environment:
      - POSTGRES_PASSWORD=postgres_password
    ports:
      - "5432:5432"
    restart: always

nginx:
  container_name: nginx
  build:
    context: ./nginx
    dockerfile: Dockerfile.dev # Change to `Dockerfile` for
production
  ports:
    - "3500:80"
  depends_on:
    - api
    - client
  restart: always

api:
  build:
    context: "./server"
    dockerfile: Dockerfile.dev # Change to `Dockerfile` for production
  volumes:
    - /app/node_modules
    - ./server:/app
  environment:
    - PGUSER=postgres
    - PGHOST=postgres
    - PGDATABASE=postgres
    - PGPASSWORD=postgres_password
    - PGPORT=5432
  depends_on:
    - postgres
  ports:
    - "8000:8000"
  restart: always

```

```
client:
  stdin_open: true
  build:
    context: ./client
    dockerfile: Dockerfile.dev # Change to `Dockerfile` for production
  volumes:
    - /app/node_modules
    - ./client:/app
  ports:
    - "5173:5173"
  depends_on:
    - api
  restart: always
```

## Explanation of the docker-compose.yml Services

### 1. Postgres:

- This service runs the PostgreSQL database.
- `environment`: Sets the environment variable `POSTGRES_PASSWORD` to `postgres_password` for the database.
- `ports`: Exposes port `5432` for database connections.
- `restart: always`: Ensures the container restarts automatically in case of a failure.

### 2. Nginx:

- Acts as a reverse proxy, routing requests to the `client` (React frontend) and `api` (Express backend) services.
- `build`: Builds the Nginx image from the `Dockerfile.dev` located in the `nginx` directory.
- `ports`: Exposes port `3500` on the host, mapping it to port `80` inside the container (where Nginx listens).

- `depends_on`: Ensures that the `api` and `client` services are started before Nginx.

### 3. API (Express):

- Runs the Express server.
- `build`: Builds the Express image from the `Dockerfile.dev` located in the `server` directory.
- `volumes`: Mounts the `server` directory and `node_modules` for hot-reloading in development.
- `environment`: Passes PostgreSQL connection environment variables to the Express server.
- `depends_on`: Ensures that the PostgreSQL service is up and running before the API service starts.
- `ports`: Exposes port `8000` for the Express API.

### 4. Client (React):

- Runs the React frontend.
- `build`: Builds the React image from the `Dockerfile.dev` located in the `client` directory.
- `volumes`: Mounts the `client` directory and `node_modules` for hot-reloading in development.
- `ports`: Exposes port `5173` for the React development server.
- `depends_on`: Ensures the API service is running before the Client starts.

## Step-by-Step Setup and Execution

### 1. Build and Start Services:

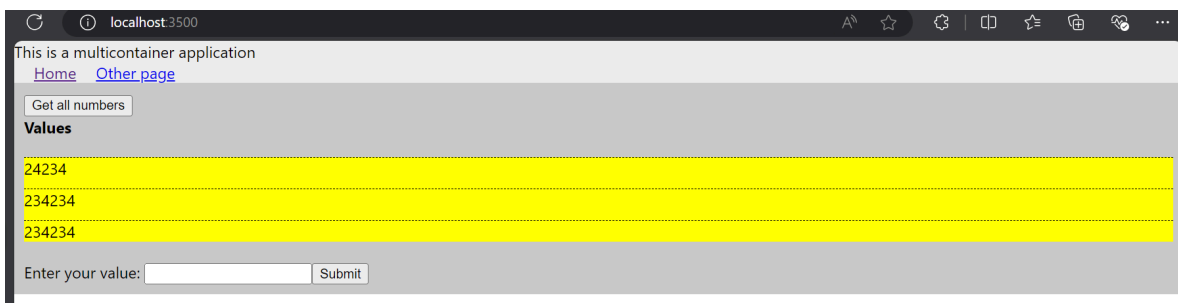
- Navigate to the root directory (`my-app/`) and run the following command to build and start all services:

```
docker-compose up --build
```

- The `--build` flag ensures that Docker builds the images before starting the containers.

## 2. Access the Application:


- Visit `http://localhost:3500` in your browser to access the frontend, which is proxied through Nginx.
- The React frontend interacts with the Express backend via the `/api` route, as defined in the Nginx configuration.



devwebpreview.png

## Production Mode

We are going to use all files in the development mode but make a couple of changes. ie our docker files will be `Docker` instead of `Docker.dev`



prod.png

## 1. ./client/nginx/default.conf

- Create Nginx folder in client folder and add `default.conf` with this code

```
server {  
    listen 5173;  
  
    location / {  
        root /usr/share/nginx/html;  
        index index.html index.html;  
        try_files $uri/ $uri/ /index.html  
    }  
}
```

## 2. ./client/Docker

- inside the client folder, add a `Docker` file

```
#use alpine Node.js runtime as a parent image and name it as builder  
FROM node:20.16.0-alpine as builder
```

```
# Set the working directory in the container
```

```
WORKDIR /app

# Copy package.json and package-lock.json
COPY package*.json ./

# Install dependencies
RUN npm i

# Copy the rest of the application code
COPY . .

RUN npm run build

FROM nginx

# Expose port 5173
EXPOSE 5173
# copy nginx configuration to the docker image
COPY ./nginx/default.conf /etc/nginx/conf.d/default.conf
# copy all builded files
COPY --from=builder /app/build /usr/share/nginx/html
```

Summary in bullet points:

- **Base Image (Builder Stage):**
  - Uses `node:20.16.0-alpine` as the base image, named as `builder`.
  - Sets the working directory to `/app` inside the container.
- **Dependency Installation:**
  - Copies `package.json` and `package-lock.json` into the container.
  - Installs Node.js dependencies using `npm i`.
- **Application Code:**
  - Copies the entire application code into the container.

- Runs the build process with `npm run build`.
- **Final Image (Nginx Stage):**
  - Uses `nginx` as the final base image.
  - Exposes port `5173` for the frontend.
  - Copies the Nginx configuration file (`default.conf`) to the appropriate location in the Nginx container.
  - Copies the built files from the `builder` stage to the Nginx web root directory (`/usr/share/nginx/html`).

### 3. ./nginx

- inside the nginx folder create this file `default.conf` add this code

```
upstream client {
    server client:5173;
}

upstream api {
    server api:8000;
}

server {
    listen 80;

    location / {
        proxy_pass http://client;
    }

    location /sockjs-node {
        proxy_pass http://client;
        proxy_http_version 1.1;
        proxy_set_header Upgrade $http_upgrade;
        proxy_set_header Connection "Upgrade";
    }
}
```

```

    location /api {
        rewrite /api/(.*) /$1 break;
        proxy_pass http://api;
    }
}

```

This configuration is for an Nginx server, typically used as a reverse proxy to direct incoming HTTP requests to different backend services based on the request's URL path. Here's a brief explanation:

#### **upstream client { ... } and upstream api { ... }**

- **upstream client { server client:5173; }**: Defines a group of servers (in this case, a single server) that Nginx will forward requests to for the **client**. The **client** service is running on port **5173**.
- **upstream api { server api:8000; }**: Similarly, this defines a group of servers for the **api**, which runs on port **8000**.

#### **server { ... }**

This block defines the configuration for handling incoming requests.

1 **listen 80;**: Configures the server to listen for HTTP requests on port **80**, the default HTTP port.

- **location / { ... }**:
  - Handles requests to the root (**/**) path.
  - **proxy\_pass http://client;**: Forwards these requests to the **client** service defined earlier, which is the application running on port **5173**.
- **location /sockjs-node { ... }**:
  - Handles WebSocket connections, specifically for **sockjs-node**.
  - **proxy\_pass http://client;**: Forwards these WebSocket requests to the **client** service.



- **proxy\_http\_version 1.1**;: Ensures HTTP/1.1 is used, which is required for WebSocket connections.
- **proxy\_set\_header Upgrade \$http\_upgrade**; and **proxy\_set\_header Connection "Upgrade"**;: These headers are necessary for handling WebSocket connections, allowing the protocol to be upgraded from HTTP to WebSocket.
- **location /api { ... }**:
  - Handles requests to the `/api` path.
  - **rewrite /api/(.\*) /\$1 break**;: This rewrites the URL by stripping the `/api` prefix before passing it to the `api` service. For example, a request to `/api/users` becomes `/users`.
  - **proxy\_pass http://api**;: Forwards the rewritten requests to the `api` service running on port `8000`.
- inside the `nginx` folder create this file `Dockerfile` add this code

```
FROM nginx
COPY ./default.conf /etc/nginx/conf.d/default.conf
```

### 3. ./server

- inside the `server` folder create this file `Dockerfile` add this code

```
#use alpine Node.js runtime as a parent image
FROM node:20.16.0-alpine

# Set the working directory in the container
WORKDIR /app

# Copy package.json and package-lock.json
COPY package*.json ./

# Install dependencies
RUN npm install
```

```
# Copy the rest of the application code
COPY . .

# Expose port 8000
EXPOSE 8000

# Command to run the application
CMD ["npm", "run", "start"]
```

#### 4. In both server and client folder

- Add `Dockerfile` with this code

```
node_modules
build
.git
.gitignore
Dockerfile
Dockerfile.dev
.dockerignore
*.log
.DS_Store
.git
```

#### 4. In the root folder create a `Docker-compose.yml` file.

```
version: "3.8"
services:
  postgres:
    image: "postgres:latest"
    environment:
      - POSTGRES_PASSWORD=postgres_password
  nginx:
    container_name: nginx
    build:
      context: ./nginx
```

```

    dockerfile: Dockerfile # Change to `Dockerfile` for production
ports:
  - "80:80"
depends_on:
  - api
  - client
restart: always
api:
  build:
    dockerfile: Dockerfile # Change to `Dockerfile` for production
    context: "./server"
  volumes:
    - /app/node_modules
    - ./server:/app
  environment:
    - PGUSER=postgres
    - PGHOST=postgres
    - PGDATABASE=postgres
    - PGPASSWORD=postgres_password
    - PGPORT=5432
client:
  stdin_open: true
  build:
    dockerfile: Dockerfile # Change to `Dockerfile` for production
    context: ./client
  volumes:
    - /app/node_modules
    - ./client:/app

```

This `docker-compose.yml` file defines a multi-service application using Docker Compose, where different containers are set up for PostgreSQL, Nginx, the API server, and the client application. Here's a breakdown of what each part does:

## Version

- **version: "3.8"**: Specifies the version of the Docker Compose file format. Version 3.8 is

compatible with the latest features in Docker.

## Services

The `services` section defines all the containers that will be started as part of this application.

### 1. postgres Service

- `image: "postgres:latest"`: Uses the latest official PostgreSQL Docker image.
- `environment:`: Sets environment variables for the PostgreSQL container.
  - `POSTGRES_PASSWORD=postgres_password`: Sets the password for the default `postgres` user in PostgreSQL.

### 2. nginx Service

- `container_name: nginx`: Names the Nginx container as `nginx`.
- `build:`: Specifies how to build the Nginx Docker image.
  - `context: ./nginx`: Uses the `nginx` directory as the build context (where Docker looks for files like the Dockerfile).
  - `dockerfile: Dockerfile`: Specifies the Dockerfile to use for building the Nginx image. It's set to `Dockerfile`, indicating that it's ready for a production environment.
- `ports:`: Maps port `80` on the host to port `80` in the container, making the Nginx server accessible via HTTP.
- `depends_on:`: Ensures that the `api` and `client` services are started before Nginx.
- `restart: always`: Ensures that the Nginx container is always restarted if it stops.

### 3. api Service

- `build:`: Specifies how to build the API server Docker image.
  - `dockerfile: Dockerfile`: Specifies the Dockerfile to use, ready for production.
  - `context: "./server"`: Uses the `server` directory as the build context.

- **volumes::** Mounts local directories into the container.
  - `/app/node_modules`: Ensures `node_modules` inside the container isn't overwritten by an empty host directory.
  - `./server:/app`: Mounts the `server` directory on the host to `/app` inside the container, allowing live development.
- **environment::** Sets environment variables to configure the API server.
  - `PGUSER=postgres`, `PGHOST=postgres`, `PGDATABASE=postgres`, `PGPASSWORD=postgres_password`, `PGPORT=5432`: These configure the connection to the PostgreSQL database running in the `postgres` service.

#### 4. client Service

- **stdin\_open: true**: Keeps the standard input open, which is often used for interactive debugging.
- **build::** Specifies how to build the client application Docker image.
  - `dockerfile: Dockerfile`: Specifies the Dockerfile for the client, set up for production.
  - `context: ./client`: Uses the `client` directory as the build context.
- **volumes::** Mounts local directories into the container.
  - `/app/node_modules`: Similar to the API service, it preserves the `node_modules` directory inside the container.
  - `./client:/app`: Mounts the `client` directory on the host to `/app` inside the container.

### Summary

This Docker Compose file sets up a full-stack application with:

- A **PostgreSQL database** (`postgres`).
- An **Nginx server** to act as a reverse proxy (`nginx`).
- An **API backend** (`api`).
- A **client frontend** (`client`).

## Build and Run using docker-compose.yml

To build and run the services defined in the `docker-compose.yml` file, you can use the following commands in your terminal:

### 1. Build the Services

```
docker-compose build
```

This command will build the Docker images for all the services (PostgreSQL, Nginx, API, and Client) as defined in the `docker-compose.yml` file. It will use the specified `Dockerfile` and context for each service.

### 2. Run the Services

```
docker-compose up
```

This command will start all the services defined in the `docker-compose.yml` file. It will also create and start any necessary networks, volumes, and dependencies.

### 3. Run the Services in Detached Mode (Optional)

If you want to run the services in the background (detached mode), you can use:

```
docker-compose up -d
```

This will start the containers in the background, allowing you to continue using the terminal for other tasks.

### 4. Stopping the Services

To stop the services that are running:

```
docker-compose down
```

This will stop and remove all the containers, networks, and volumes created by `docker-compose up`.

## Pushing local Docker images to DockerHub Repositories

- sign in to <https://hub.docker.com/> and create a repository called `react-app-prod`, `express-app-prod` and `nginx-prod`
- Go back to your project and cd into the root of `client`
- Log in to the Docker registry:

```
docker login
```

- Enter your username and password when prompted. If you are using Docker Hub, these would be your Docker Hub credentials. For private registries, use the appropriate credentials.

## Client(React-app)

- 
- To push an image use

```
docker push my-dockerhub-username/my-repo:tag
```

- For this example I'll build the client image with `my-dockerhub-username/my-repo:tag`

```
docker build -t kelchospense/react-app-prod:v1 .
```

- Then push

```
docker push kelchospense/react-app-prod:v1
```

## Server(Express-app)

- To build server go to the `./server` dir

```
docker build -t kelchospense/express-app-prod:v1 .
```

- Then push

```
docker push kelchospense/express-app-prod:v1
```



# Restore the old Context Menu in Windows 11

1. Right-click the Start button and choose Windows Terminal.
2. Copy the command from below, paste it into Windows Terminal Window, and press enter.
3. 

```
reg.exe add "HKCU\Software\Classes\CLSID\{86ca1aa0-34aa-4e8b-a509-50c905bae2a2}\InprocServer32" /f /ve
```
4. Restart File Explorer or your computer for the changes to take effect.
5. You would see the Legacy Right Click Context menu by default.

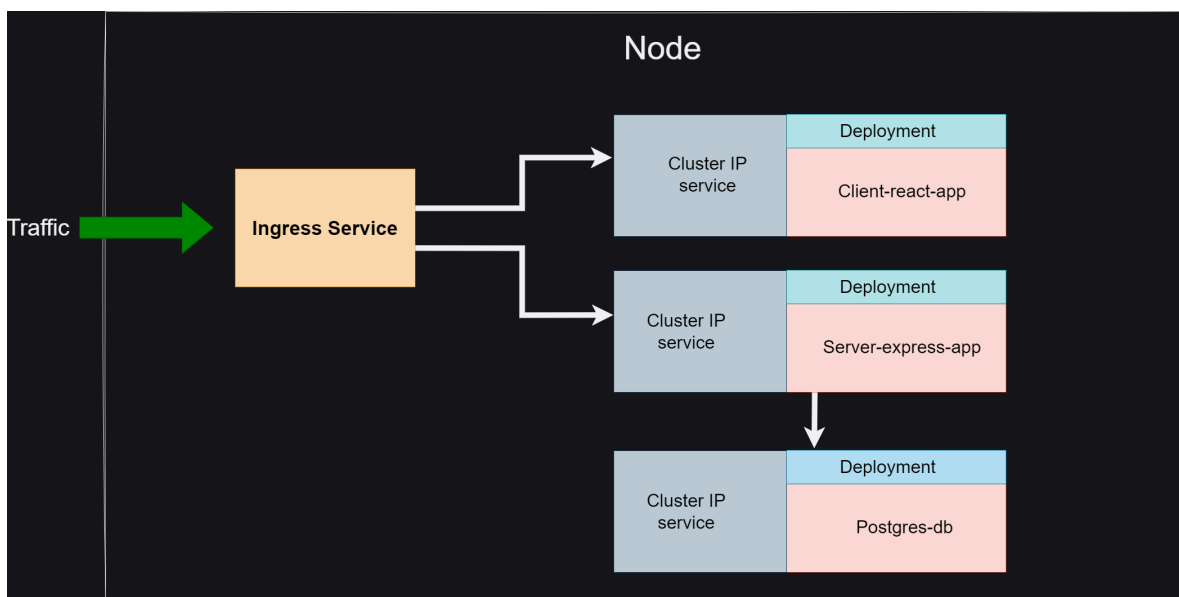
## Restore Modern Context menus in Windows 11

- To undo this change, in a Terminal Window, execute this command:

```
reg.exe delete "HKCU\Software\Classes\CLSID\{86ca1aa0-34aa-4e8b-a509-50c905bae2a2}" /f
```

# Kubernetes Multi Container Deployment | React | Node.js | Postgres | Ingress Nginx | step by step

## Chapter 1: Introduction



kubernetes.png

This diagram illustrates a Kubernetes deployment architecture for a full-stack application consisting of a React frontend, a Node.js (Express) backend, and a PostgreSQL database, with traffic management handled by an Ingress service using Nginx. Here's a step-by-step explanation of how this setup works:

### 1. Traffic Ingress via Ingress Service

- **Ingress Service:**
  - The first component of this architecture is the Ingress service, which acts as the entry point for all external traffic into the Kubernetes cluster.

- In this context, Nginx is likely used as the ingress controller to manage incoming HTTP/HTTPS requests. It routes the traffic based on defined rules to the appropriate services inside the cluster.

## 2. Routing to Cluster IP Services

- **Cluster IP Services:**
  - These services act as internal load balancers within the Kubernetes cluster, distributing incoming requests from the Ingress to the appropriate pods.
  - There are two Cluster IP services in this diagram:
    - One for the **Client-react-app**.
    - One for the **Server-express-app**.

## 3. Client-React-App Deployment

- **Deployment:**
  - This is where the React frontend is deployed.
  - The React application is served as static files and is accessed by the user's browser.
  - It communicates with the backend (Server-express-app) for dynamic data.
  - Traffic destined for the React app is routed by the Ingress service to this deployment via the corresponding Cluster IP service.

## 4. Server-Express-App Deployment

- **Deployment:**
  - This is where the backend server, built with Node.js and Express, is deployed.
  - The server handles API requests from the frontend and manages business logic. It also interacts with the PostgreSQL database to fetch or store data.
  - The Ingress service routes traffic intended for the backend to this deployment through its associated Cluster IP service.

## 5. Postgres-DB Deployment

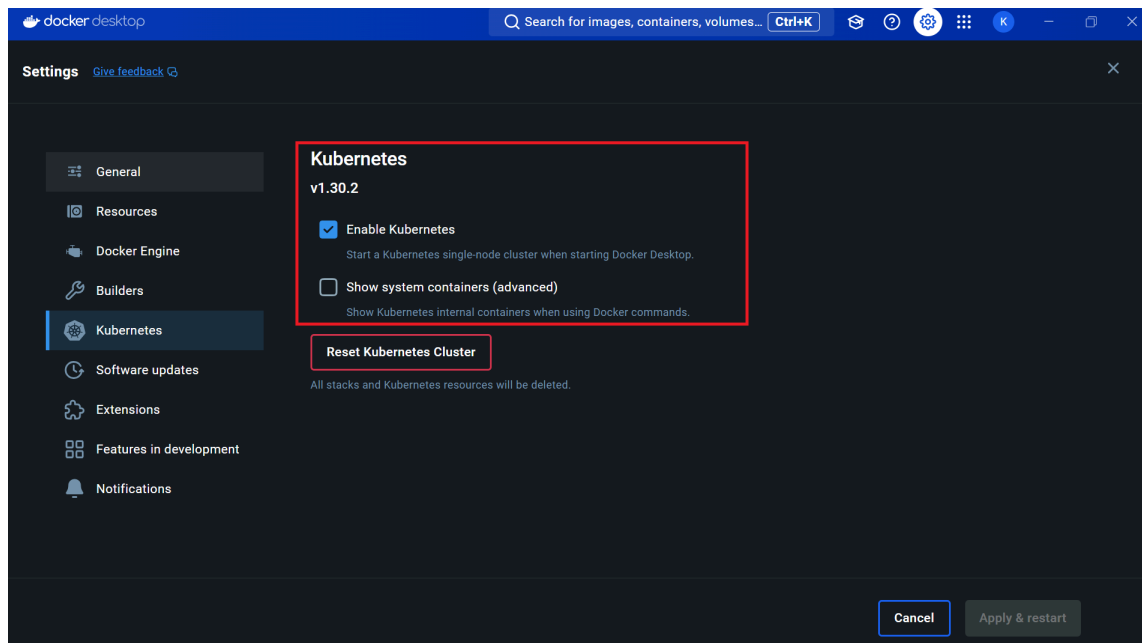
- **Deployment:**
  - The PostgreSQL database is deployed here.
  - It stores all persistent data for the application and is accessed by the Server-express-app for CRUD operations.
  - This deployment is not exposed to external traffic but is accessed internally by the Server-express-app.

## 6. Inter-Deployment Communication

- **Server-Express-App to Postgres-DB:**
  - The backend server (Node.js Express app) communicates with the PostgreSQL database deployment internally within the Kubernetes cluster.
  - The communication is facilitated via internal DNS or direct service names in Kubernetes, ensuring that the database is secure and only accessible by the backend service.

## Install Kubernetes via Docker

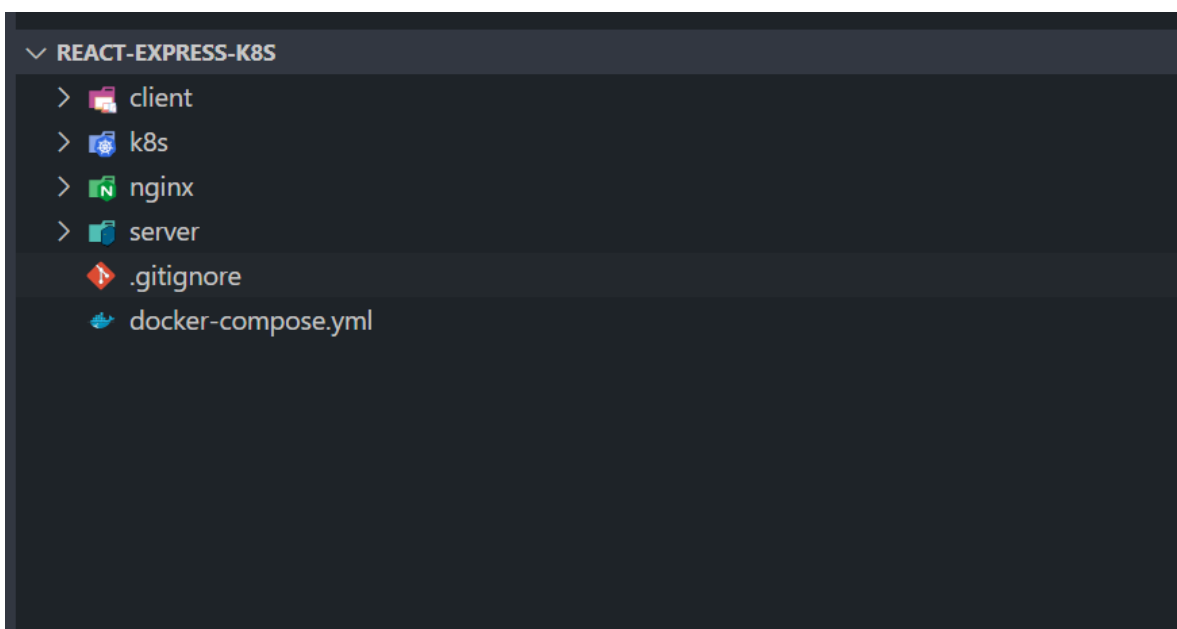
- Please make sure this option is checked below.



Screenshot 2024-08-13 170438.png

- Run this on your terminal `kubectl version` this command is used to check the version of the kubectl client and the version of the Kubernetes server (i.e., the Kubernetes control plane) that your kubectl is communicating with.

## Create K8S folder



ks8-folder-structure.png

## 1. Add client-deployment.yml file with this code

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: client-deployment
spec:
  replicas: 1
  selector:
    matchLabels:
      component: web
  template:
    metadata:
      labels:
        component: web
    spec:
      containers:
        - name: client
          image: kelchospense/react-app-prod
          ports:
            - containerPort: 5173
```

This YAML configuration is a Kubernetes manifest that defines a **Deployment** for a React application. Here's a brief explanation of each section:

- **apiVersion: apps/v1**

Specifies the version of the Kubernetes API being used. The `apps/v1` version is used for managing deployments, stateful sets, and other application-related resources.

## 2. kind: Deployment

- Indicates that this manifest describes a **Deployment** resource. A Deployment is responsible for managing a set of identical pods, ensuring that the desired number of pods are running, and handling updates to the application.

## 3. metadata

- **name: client-deployment:** Assigns a name to the Deployment. In this case, it's called `client-deployment`.

#### 4. spec

- The `spec` section defines the desired state of the Deployment.
- **replicas: 1:** Specifies that only one replica (pod) of this application should be running. This can be scaled up if needed.
- **selector:**
  - **matchLabels:** This section defines a label selector. It ensures that the Deployment manages only the pods that have the label `component: web`.
- **template:**
  - Defines the template for the pods that will be created by this Deployment.
  - **metadata:**
    - **labels:** Labels added to the pod. Here, the label `component: web` is assigned to the pod. This label must match the `matchLabels` selector in the Deployment spec.
  - **spec:**
    - **containers:** Defines the container(s) that will run inside the pod.
      - **name: client:** The name of the container running in the pod.
      - **image: kelchospense/react-app-prod:** The Docker image that the container will run. In this case, it's a React application that has been containerized and pushed to a Docker registry under the name `kelchospense/react-app-prod`.
      - **ports:**
        - **containerPort: 5173:** Specifies the port on which the application inside the container listens. This port will be exposed by the container so that it can be accessed by other services or users.

#### Testing the `client-deployment.yml`

- Go to the root dir and open terminal.

```
kubectl apply -f k8s
```

- Check the pod is running

```
kubectl get pods
```

- should get

```
client-deployment-5656f88b7c-2p9cs    1/1    Running    2
18m
```



# What is Kubernetes?

Kubernetes, often abbreviated as "K8s," is an open-source platform designed to automate the deployment, scaling, and management of containerized applications. It's crucial in modern cloud-native application development for several reasons:

## 1. Scalability

- Kubernetes allows you to easily scale your application up or down based on demand. This ensures that your application can handle traffic spikes without manual intervention.

## 2. High Availability

- Kubernetes ensures that your application is always up and running by automatically restarting failed containers, replicating services across nodes, and managing rolling updates.

## 3. Efficient Resource Utilization

- It schedules containers based on available resources and requirements, optimizing the use of underlying infrastructure. This ensures cost-effective usage of computing resources.

## 4. Self-Healing

- Kubernetes automatically monitors the health of your applications and containers. If a container fails, it restarts it. If a node fails, it redistributes the workloads.

## 5. Automated Rollouts and Rollbacks

- You can define how your application should be updated. Kubernetes will gradually roll out changes while monitoring the health of your application, allowing for automated rollbacks if something goes wrong.

## 6. Portability and Flexibility

- Kubernetes abstracts away the underlying infrastructure, making it easier to move workloads between different environments (on-premises, public cloud, hybrid cloud).

### Use Cases:

- **Microservices Architecture:** Kubernetes is ideal for deploying microservices because it can manage a large number of interdependent services with ease.
- **DevOps Pipelines:** Automate deployment, scaling, and operations of application containers across clusters of hosts.
- **Hybrid Cloud Deployments:** Consistently deploy applications across on-premises and public cloud environments.
- **Batch Processing and CI/CD:** Efficiently run batch jobs, continuous integration, and continuous deployment tasks with Kubernetes.

## Practical Example: Kubernetes on Docker Desktop

Let's create a simple example where we deploy a Node.js application on Kubernetes using Docker Desktop.

### Step 0: Set Up Kubernetes on Docker Desktop

#### 1. Install Docker Desktop:

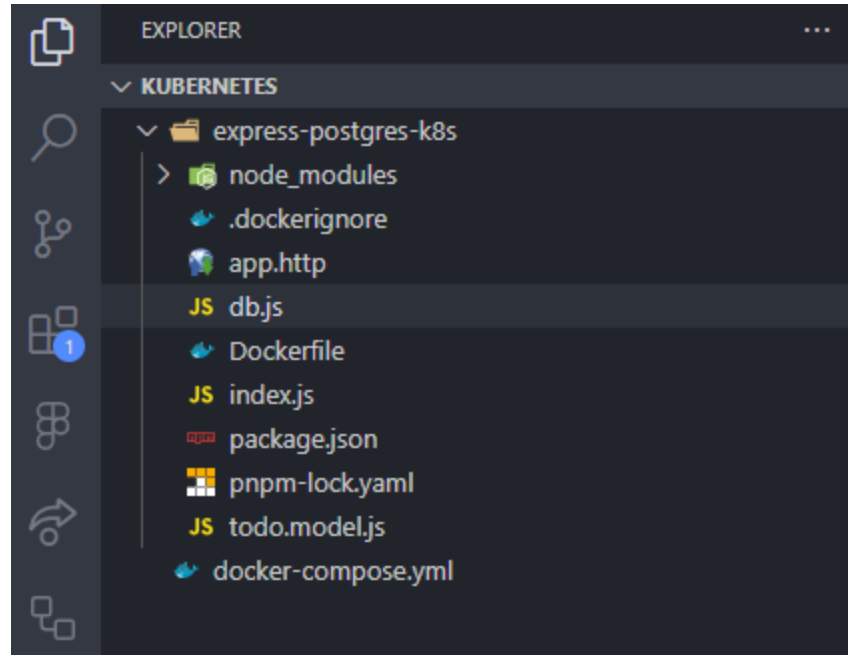
- Download and install Docker Desktop from the official site (<https://www.docker.com/products/docker-desktop>).

#### 2. Enable Kubernetes:

- Open Docker Desktop, go to **Settings > Kubernetes** and check "Enable Kubernetes".

- Click "Apply & Restart". Docker Desktop will install and start a local Kubernetes cluster.

## Step 1: Set Up the Node.js Application



express-k8s-folder-structure.png

### 1. Initialize the Node.js Project:

```
mkdir express-postgres-k8s
cd express-postgres-k8s
npm init -y
npm install express pg sequelize
```

- **express**: Framework to build the API.
- **pg**: PostgreSQL client for Node.js.
- **sequelize**: ORM (Object-Relational Mapping) for PostgreSQL.

### 2. Set Up Sequelize and PostgreSQL Connection:

Create a file `db.js` to configure Sequelize and connect to PostgreSQL:

```

const { Sequelize } = require('sequelize');

// Create a new Sequelize instance
const sequelize = new Sequelize('todo_db', 'postgres', 'password', {
  host: 'localhost',
  dialect: 'postgres',
});

// Test the connection
sequelize
  .authenticate()
  .then(() => {
    console.log('Connection to PostgreSQL has been established successfully.');
```

```

  })
  .catch(err => {
    console.error('Unable to connect to the database:', err);
  });

module.exports = sequelize;

```

- Database Name: `todo_db`
- User: `postgres`
- Password: `password`
- Host: `localhost` (will be updated for Docker/Kubernetes)

### 3. Define the To-Do Model:

Create a file `todo.model.js` to define the To-Do model:

```

const { DataTypes } = require('sequelize');
const sequelize = require('./db');

const Todo = sequelize.define('Todo', {
  id: {
    type: DataTypes.INTEGER,
```

```

    autoIncrement: true,
    primaryKey: true,
  },
  title: {
    type: DataTypes.STRING,
    allowNull: false,
  },
  completed: {
    type: DataTypes.BOOLEAN,
    defaultValue: false,
  },
});

// Sync the model with the database
sequelize.sync()
  .then(() => {
    console.log('Todo model synced with the database.');
```

```

  })
  .catch(err => {
    console.error('Failed to sync model with the database:', err);
  });

module.exports = Todo;

```

#### 4. Create the Express Server with CRUD Routes:

Create a file `index.js` to define the Express server and CRUD routes:

```

const express = require('express');
const Todo = require('./todo.model');

const app = express();
app.use(express.json());

// Create a new to-do
app.post('/todos', async (req, res) => {
  try {

```

```

    const todo = await Todo.create(req.body);
    res.status(201).json(todo);
  } catch (err) {
    res.status(400).json({ error: err.message });
  }
});

// Get all to-dos
app.get('/todos', async (req, res) => {
  try {
    const todos = await Todo.findAll();
    res.status(200).json(todos);
  } catch (err) {
    res.status(400).json({ error: err.message });
  }
});

// Get a single to-do by ID
app.get('/todos/:id', async (req, res) => {
  try {
    const todo = await Todo.findByPk(req.params.id);
    if (todo) {
      res.status(200).json(todo);
    } else {
      res.status(404).json({ error: 'To-do not found' });
    }
  } catch (err) {
    res.status(400).json({ error: err.message });
  }
});

// Update a to-do
app.put('/todos/:id', async (req, res) => {
  try {
    const todo = await Todo.findByPk(req.params.id);
    if (todo) {
      await todo.update(req.body);
      res.status(200).json(todo);
    }
  }
});

```

```

    } else {
      res.status(404).json({ error: 'To-do not found' });
    }
  } catch (err) {
    res.status(400).json({ error: err.message });
  }
});

// Delete a to-do
app.delete('/todos/:id', async (req, res) => {
  try {
    const todo = await Todo.findByPk(req.params.id);
    if (todo) {
      await todo.destroy();
      res.status(204).end();
    } else {
      res.status(404).json({ error: 'To-do not found' });
    }
  } catch (err) {
    res.status(400).json({ error: err.message });
  }
});

// Start the server
const PORT = process.env.PORT || 3000;
app.listen(PORT, () => {
  console.log(`Server is running on port ${PORT}`);
});

```

## 5. Create a PostgreSQL Database:

Before testing, create the database `todo_db` in PostgreSQL:

```

psql -U postgres
CREATE DATABASE todo_db;

```

## 6. Test the Application Locally:

Run the Node.js application:

```
node index.js
```

- Use Postman or cURL to test the CRUD operations against `http://localhost:3000/todos`.

## Step 2: Dockerize the Application

### 1. Create a Dockerfile:

```
FROM node:14

WORKDIR /app

COPY package*.json ./
RUN npm install

COPY . .

EXPOSE 3000

CMD ["node", "index.js"]
```

### 2. Create a `docker-compose.yml` outside `express-postgres-k8s` folder for Local Development:

We'll use Docker Compose to run the application along with PostgreSQL.

```
version: '3.8'

services:
  postgres:
    image: postgres:13
    environment:
      POSTGRES_USER: postgres
      POSTGRES_PASSWORD: password
      POSTGRES_DB: todo_db
```



```

    ports:
      - "5432:5432"
    volumes:
      - postgres-data:/var/lib/postgresql/data

  app:
    build:
      dockerfile: Dockerfile
      context: ./express-postgres-k8s
    environment:
      DB_HOST: postgres
      DB_USER: postgres
      DB_PASSWORD: password
      DB_NAME: todo_db
    ports:
      - "3000:3000"
    depends_on:
      - postgres

  volumes:
    postgres-data:

```

- **Explanation:**

- This `docker-compose.yml` sets up two services: `postgres` and `app`.
- The Node.js app depends on the PostgreSQL service and uses environment variables to connect to the database.

## 2. Update `db.js` to Use Environment Variables:

Update the `db.js` file to use the environment variables defined in `docker-compose.yml`:

```

const { Sequelize } = require('sequelize');

const sequelize = new Sequelize(process.env.DB_NAME,
process.env.DB_USER, process.env.DB_PASSWORD, {

```

```

    host: process.env.DB_HOST,
    dialect: 'postgres',
  });

  sequelize.authenticate()
    .then(() => {
      console.log('Connection to PostgreSQL has been established successfully.');
```

```
    })
```

```
    .catch(err => {
```

```
      console.error('Unable to connect to the database:', err);
```

```
    });
```

```
module.exports = sequelize;
```

3. **Testing the API:** install REST Client extension and create `app.http` and add this code

### Create a new To-Do

POST http://localhost:3000/todos

Content-Type: application/json

```
{
  "title": "Buy groceries"
}
```

### Get all To-Dos

GET http://localhost:3000/todos

### Get a single To-Do by ID

GET http://localhost:3000/todos/1

### Update a To-Do (mark as completed)

PUT http://localhost:3000/todos/1

Content-Type: application/json

```
{
```

```
"completed": true
}
```

### Delete a To-Do

```
DELETE http://localhost:3000/todos/1
```

## 5. Build and Run the Docker Compose Setup:

```
docker-compose up --build
```

- This command will build the Node.js Docker image, start the PostgreSQL service, and run the application.
- The app will be accessible at `http://localhost:3000`.

## Step 3: Deploy to Kubernetes

Now, let's deploy the application on Kubernetes.

### 1. Push the Docker Image to Docker Hub:

```
docker build -t <your-dockerhub-username>/node-postgres-todo:v1 .
docker push <your-dockerhub-username>/node-postgres-todo:v1
```

### 2. Create Kubernetes Manifests:

#### 1. PostgreSQL Deployment (postgres-deployment.yaml):

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: postgres-deployment
spec:
  replicas: 1
  selector:
    matchLabels:
      app: postgres
```

```

template:
  metadata:
    labels:
      app: postgres
  spec:
    containers:
      - name: postgres
        image: postgres:13
        ports:
          - containerPort: 5432
        env:
          - name: POSTGRES_USER
            value: "postgres"
          - name: POSTGRES_PASSWORD
            value: "password"
          - name: POSTGRES_DB
            value: "todo_db"
        volumeMounts:
          - mountPath: /var/lib/postgresql/data
            name: postgres-storage
    volumes:
      - name: postgres-storage
        persistentVolumeClaim:
          claimName: postgres-pvc

```

## 2. PostgreSQL Persistent Volume Claim (postgres-pvc.yaml):

```

apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: postgres-pvc
spec:
  accessModes:
    - ReadWriteOnce
  resources:

```

```
requests:
  storage: 1Gi
```

### 3. PostgreSQL Service (postgres-service.yaml):

```
apiVersion: v1
kind: Service
metadata:
  name: postgres-service
spec:
  selector:
    app: postgres
  ports:
    - protocol: TCP
      port: 5432
  type: ClusterIP
```

### 4. Node.js Application Deployment (node-app-deployment.yaml):

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: node-app-deployment
spec:
  replicas: 3
  selector:
    matchLabels:
      app: node-app
  template:
    metadata:
      labels:
        app: node-app
    spec:
      containers:
        - name: node-app
          image: <your-dockerhub-username>/node-postgres-todo:v1
```

```

ports:
- containerPort: 3000
env:
- name: DB_HOST
  value: "postgres-service"
- name: DB_USER
  value: "postgres"
- name: DB_PASSWORD
  value: "password"
- name: DB_NAME
  value: "todo_db"

```

### 5. Node.js Application Service (node-app-service.yaml):

```

apiVersion: v1
kind: Service
metadata:
  name: node-app-service
spec:
  selector:
    app: node-app
  ports:
    - protocol: TCP
      port: 80
      targetPort: 3000
  type: LoadBalancer

```

### 3. Deploy to Kubernetes:

```

kubectl apply -f postgres-pvc.yaml
kubectl apply -f postgres-deployment.yaml
kubectl apply -f postgres-service.yaml
kubectl apply -f node-app-deployment.yaml
kubectl apply -f node-app-service.yaml

```

### 4. Verify the Deployment:

- Check the pods and services:

```
kubectl get pods  
kubectl get svc
```

- Once the services are up and running, the Node.js app will be accessible via the external IP of the `node-app-service`.

## Conclusion

You've now created a full-stack Node.js application using Express and PostgreSQL with a simple To-Do list example. You've containerized the application using Docker, orchestrated it using Docker Compose for local development, and finally deployed it on Kubernetes. This example demonstrates how Kubernetes can help manage a multi-service application, providing scalability, load balancing, and resilience.

# Kubernetes Cheat Sheet

Here's a Kubernetes cheat sheet with some commonly used commands, including how to delete pods and more:

## 1. Basic Commands

- Get cluster information:

```
kubectl cluster-info
```

- View nodes in the cluster:

```
kubectl get nodes
```

- View all resources (pods, services, etc.):

```
kubectl get all
```

## 2. Pods

- List all pods in the current namespace:

```
kubectl get pods
```

- List pods in all namespaces:

```
kubectl get pods --all-namespaces
```

- Get detailed information about a pod:

```
kubectl describe pod <pod_name>
```

- Delete a specific pod:



```
kubectl delete pod <pod_name>
```

- **Force delete a pod (without waiting for graceful termination):**

```
kubectl delete pod <pod_name> --grace-period=0 --force
```

- **Get logs of a specific pod:**

```
kubectl logs <pod_name>
```

- **Stream logs of a specific pod:**

```
kubectl logs -f <pod_name>
```

### 3. Deployments

- **List all deployments:**

```
kubectl get deployments
```

- **Describe a specific deployment:**

```
kubectl describe deployment <deployment_name>
```

- **Create a deployment:**

```
kubectl create deployment <deployment_name> --image=<image_name>
```

- **Update a deployment with a new image:**

```
kubectl set image deployment/<deployment_name> <container_name>=  
<new_image_name>
```

- **Scale a deployment:**

```
kubectl scale deployment <deployment_name> --replicas=  
<number_of_replicas>
```

- Delete a deployment:

```
kubectl delete deployment <deployment_name>
```

## 4. Services

- List all services:

```
kubectl get services
```

- Describe a specific service:

```
kubectl describe service <service_name>
```

- Expose a deployment as a service:

```
kubectl expose deployment <deployment_name> --type=<type> --port=  
<port> --target-port=<target_port>
```

- Delete a service:

```
kubectl delete service <service_name>
```

## 5. Namespaces

- List all namespaces:

```
kubectl get namespaces
```

- Create a new namespace:

```
kubectl create namespace <namespace_name>
```

- **Delete a namespace:**

```
kubectl delete namespace <namespace_name>
```

- **Switch to a different namespace:**

```
kubectl config set-context --current --namespace=<namespace_name>
```

## 6. ConfigMaps & Secrets

- **Create a ConfigMap:**

```
kubectl create configmap <configmap_name> --from-literal=<key>=<value>
```

- **Create a Secret:**

```
kubectl create secret generic <secret_name> --from-literal=<key>=  
<value>
```

- **Describe a ConfigMap or Secret:**

```
kubectl describe configmap <configmap_name>  
kubectl describe secret <secret_name>
```

- **Delete a ConfigMap or Secret:**

```
kubectl delete configmap <configmap_name>  
kubectl delete secret <secret_name>
```

## 7. Others

- **Apply a configuration from a file:**

```
kubectl apply -f <file_name>.yaml
```

- **Delete resources from a file:**

```
kubectl delete -f <file_name>.yaml
```

- **View events:**

```
kubectl get events
```

- **View the current context:**

```
kubectl config current-context
```

This cheat sheet covers the most common operations in Kubernetes. It's helpful to keep this handy when working with Kubernetes clusters!